AI SECURITY STANDARD

Document ID: DOC-024 | Version: 1.0 | Effective: 2026-01-01 | Status: ACTIVE | Type: Standard

1. Purpose

Define mandatory security controls for enterprise use of AI systems.

2. Data classification

AI processing must respect the enterprise information classification. Confidential or restricted information may be processed only in an environment explicitly approved for that classification and use case.

3. Minimum controls

Control	Requirement
Identity and access	Use enterprise-approved authentication and least-privilege access where supported.
Data handling	Input only data permitted for the approved tool, use case, and classification.
Provider terms	Approved processing and retention terms must be established where external providers process enterprise data.
Logging and traceability	Maintain usage or audit evidence when required by the risk classification or use case.
Secrets	Credentials, private keys, passwords, and equivalent authentication secrets must not be entered into general AI prompts.
Incident handling	Suspected exposure through an AI system must follow the Security Incident SOP.

4. Contract analysis

Supplier contracts may be processed only in an approved AI environment authorized for the contract's information classification. Users must not assume that a consumer-accessible AI tool is approved for confidential contract content.

5. Output handling

AI output may contain inaccurate, incomplete, or unsupported content. Security controls do not make model output authoritative; human review remains required according to the business use case.

Related documents

AI Governance Policy; Information Security Policy; Security Incident SOP.